

DDA5001: Homework #1

Due on September 28, 2025 at 23:59

Professor DAI Zhongxiang Term 1, 2025-2026

LIU Yibo

Problem 1. Concept and Fundamental Knowledge

Solution:

(a) The main difference between supervised learning and unsupervised learning is **the use of labeled data**. In supervised learning, each data point is tagged with a correct output, while in unsupervised learning, there are no predefined output labels. The character of the data they utilize determines the tasks they commonly deal with: supervised Learning used to solve *classification* and *regression*; unsupervised learning usually suits the conditions of *Clustering*, *association*, and *dimensionality reduction*.

(b) 1) **False**. Regression is used to predict a continuous value, such as stock prices, while classification fits categorical labels.

2) **True**.

3) **Flase**. For any linearly separable dataset, there are typically infinite possible hyperplanes that can separate the data.

4) **False**. Least squares is only equivalent to an MLE under the assumption that the errors are independent and Gaussian distributed with a mean of zero. Without the Gaussian error assumption, it is not an MLE.

(c) Denote the columns of full rank matrix \mathbf{X} as $\mathbf{X}_1, \mathbf{X}_2, \dots, \mathbf{X}_n$, so that

$$\mathbf{X} = [\mathbf{X}_1, \mathbf{X}_2, \dots, \mathbf{X}_n]. \quad (1)$$

With the independence of the columns, for any vector $\mathbf{v} = [v_1, v_2, \dots, v_n]^T \neq \mathbf{0}$, we have

$$\mathbf{y} = \mathbf{X}\mathbf{v} = \sum_{i=1}^n \mathbf{X}_i v_i \neq \mathbf{0}. \quad (2)$$

Therefore, we obtain

$$\mathbf{v}^T (\mathbf{X}^T \mathbf{X}) \mathbf{v} = (\mathbf{X}\mathbf{v})^T (\mathbf{X}\mathbf{v}) = \mathbf{y}^T \mathbf{y} > 0, \quad (3)$$

which says that $\mathbf{X}^T \mathbf{X}$ is positive definite.

Problem 2. Least Squares Without Full Column Rank

Solution:

(a) With the singular value decomposition (SVD), we can substitute $X = V\Sigma_1 U_1^T$ into the original form:

$$\|X\theta - y\|_2^2 = \|V\Sigma_1 U_1 \theta - y\|_2^2, \quad (4)$$

where $V \in \mathbb{R}^{n \times n}$ is orthogonal, $\Sigma_1 \in \mathbb{R}^{n \times n}$ is diagonal with positive singular value $\sigma_1 \geq \sigma_2 \geq \dots \geq \sigma_n$, and $U_1^T \in \mathbb{R}^{n \times d}$ is a semi-orthogonal matrix. Let $A := V\Sigma_1$ and $z := U_1^T \theta$, the initial problem is equivalent to

$$\min_z \|Az - y\|_2^2. \quad (5)$$

The minimal solutions can be obtained when $Az^* = y$, i.e., $z^* = A^{-1}y = \Sigma_1^{-1}V^T y$. Let $U = [U_1, U_2]$ and $w := U_2^T \theta$, then we have

$$\begin{aligned} \theta &= U(U^T \theta) = [U_1, U_2] \begin{bmatrix} U_1^T \theta \\ U_2^T \theta \end{bmatrix} \\ &= U_1(U_1^T \theta) + U_2(U_2^T \theta) \\ &= U_1 z + U_2 w. \end{aligned} \quad (6)$$

Here, w is an arbitrary vector in \mathbb{R}^{n-d} , leading to the infinite solutions of least squares without full rank. This is because $n < d$, thus the null space of X is non-trivial (Since $XU_2 = (V\Sigma_1 U_1^T)U_2 = V\Sigma_1(U_1^T U_2) = 0$, the column of U_2 form an orthonormal basis for the null space of X). Any vector from the null space of X can be added to a particular solution $\hat{\theta}$ without changing the result of $X\theta$.

(b) Let the objective function be $J(\theta)$:

$$\begin{aligned} J(\theta) &= \|X\theta - y\|_2^2 + \lambda \|\theta\|_2^2 \\ &= (\theta^T X^T - y^T)(X\theta - y) + \lambda \theta^T \theta \\ &= \theta^T X^T X \theta - 2y^T X \theta + y^T y + \lambda \theta^T \theta. \end{aligned} \quad (7)$$

Then we take the gradient of $J(\theta)$:

$$\nabla_{\theta} J(\theta) = 2X^T X \theta - 2X^T y + 2\lambda \theta. \quad (8)$$

Solving for θ by setting the gradient to zero, we obtain:

$$(X^T X + \lambda I)\theta = X^T y. \quad (9)$$

Since $X^T X$ is a positive semi-definite matrix and λI is a positive matrix, $(X^T X + \lambda I)$ is invertible. Thus, we have

$$\theta = (X^T X + \lambda I)^{-1} X^T y. \quad (10)$$

Problem 3. Robust Linear Regression

Solution:

(a) Denote the probability of observing the given data y for a specific parameter θ is $L(\theta)$, $L(\theta) = p(y|X, \theta)$. Since the error term $\epsilon \stackrel{i.i.d}{\sim} \mathcal{N}(0, \Sigma)$, y_i are also conditionally independent. Therefore, we obtain the total likelihood:

$$L(\theta) = \prod_{i=1}^n p(y_i|x_i, \theta), \quad (11)$$

Given an input x and θ , the randomness of y only comes from the error term ϵ . Thus, we have

$$p(y_i|x_i, \theta) = p_\epsilon(\epsilon_i) = p_\epsilon(y_i - x_i^T \theta) = \frac{1}{2b} \exp\left(-\frac{|y_i - x_i^T \theta|}{b}\right). \quad (12)$$

Substituting $p(y_i|x_i, \theta)$ into Eq.(11), the log-likelihood $\mathcal{L}(\theta)$ can be derived as:

$$\begin{aligned} \mathcal{L}(\theta) &= \log L(\theta) \\ &= \log \left(\prod_{i=1}^n \frac{1}{2b} \exp\left(-\frac{|y_i - x_i^T \theta|}{b}\right) \right) \\ &= \sum_{i=1}^n \log \frac{1}{2b} \exp\left(-\frac{|y_i - x_i^T \theta|}{b}\right) \\ &= -n \log 2b + \sum_{i=1}^n -\frac{|y_i - x_i^T \theta|}{b}. \end{aligned} \quad (13)$$

The maximum likelihood estimation θ^* is the value of θ that maximizes the log-likelihood function:

$$\begin{aligned} \theta^* &= \arg \max_{\theta} \mathcal{L}(\theta) \\ &= \arg \max_{\theta} \left(-n \log 2b - \frac{1}{b} \sum_{i=1}^n |y_i - x_i^T \theta| \right) \\ &= \arg \min_{\theta} \left(n \log 2b + \frac{1}{b} \sum_{i=1}^n |y_i - x_i^T \theta| \right) \\ &= \arg \min_{\theta} \sum_{i=1}^n |y_i - x_i^T \theta|. \end{aligned} \quad (14)$$

This equation represents the L1-norm of the residual vector $y - X\theta$. Therefore, when ϵ follows the Laplace distribution, the machine learning problem formulation for estimation θ^* is

$$\min_{\theta} \|X\theta - y\|_1. \quad (15)$$

(b) We first define the residual vector as $\mathbf{r} = X\theta - y$, with components $r_j = x_j^T \theta - y_j, j = 1, \dots, n$, where x_j is the j -th row of X .

The loss function is defined as:

$$\mathcal{L}(\theta) = H_{\mu}(\mathbf{r}) = \sum_{j=1}^n h_{\mu}(r_j) \quad (16)$$

Then, the derivative of the Huber function is:

$$h'_{\mu}(z) = \begin{cases} \frac{z}{\mu}, & \text{if } |z| \leq \mu \\ \text{sgn}(z), & \text{if } |z| > \mu \end{cases} \quad (17)$$

Here, sgn is a function that meets

$$\text{sgn}(z) = \begin{cases} -1, & \text{if } z < 0 \\ 0, & \text{if } z = 0 \\ 1, & \text{if } z > 0 \end{cases} \quad (18)$$

This derivative is continuous and well-defined everywhere.

The j -th component of the gradient vector $\nabla_r \mathcal{L}(\theta)$ is $h'_\mu(r_j)$. Define a vector $g = \nabla_r H_\mu(r)$ such that its components are given by:

$$g_j = h'_\mu(r_j) = \begin{cases} r_j/\mu, & \text{if } |r_j| \leq \mu \\ \text{sgn}(r_j), & \text{if } |r_j| > \mu \end{cases} \quad (19)$$

where $r_j = x_j^T \theta - y_j$.

By the vector chain rule, we can obtain the gradient is

$$\nabla_\theta \mathcal{L}(\theta) = \left(\frac{\partial r}{\partial \theta} \right)^T \nabla_r H_\mu(r) \quad (20)$$

Since $r = X\theta - y$, the Jacobian matrix $\frac{\partial r}{\partial \theta}$ is simply X . With the $g = \nabla_r H_\mu(r)$, we have

$$\nabla_\theta \mathcal{L}(\theta) = X^T g. \quad (21)$$

(c) The implementation process is in the code file p3.

The results of programming show that: As expected, the error decreases rapidly as the number of iterations increases, indicating that the gradient descent algorithm is effectively minimizing the Huber loss function and that our estimated parameters θ are converging to the true values θ^* .




Figure 1: Error Convergence of L1 Estimator via Huber Smoothing

Problem 4. Convergence of The Perceptron for Linearly Separable Data

Solution:

(a) Since the data is linearly separable, for every i , we have

$$y_i(\theta^{*\top} x_i) > 0. \quad (22)$$

we can obtain that θ^* correctly classifies all samples. Because n is the number of training samples, the set $\{y_i(\theta^{*\top} x_i) \mid 1 \leq i \leq n\}$ is a finite collection of positive numbers. The minimum of a finite set of positive numbers is positive. Thus, we can obtain that

$$\rho = \min_{1 \leq i \leq n} y_i(\theta^{*\top} x_i) > 0. \quad (23)$$

(b) According to the perceptron update rule $\theta_k = \theta_{k-1} + y_{k-1}x_{k-1}$, taking the inner product with θ^* , we can obtain:

$$\begin{aligned} \theta_k^\top \theta^* &= (\theta_{k-1} + y_{k-1}x_{k-1})^\top \theta^* \\ &= \theta_{k-1}^\top \theta^* + y_{k-1}(x_{k-1}^\top \theta^*). \end{aligned} \quad (24)$$

Since (x_{k-1}, y_{k-1}) is misclassified by θ_{k-1} , it is still correctly classified by θ^* according to the linear separability. Combining with the definition of ρ , we obtain:

$$y_{k-1}(\theta^{*\top} x_{k-1}) \geq \rho. \quad (25)$$

Substituting it into the above formula, we obtain

$$\theta_k^\top \theta^* \geq \theta_{k-1}^\top \theta^* + \rho. \quad (26)$$

Next, we prove $\theta_k^\top \theta^* \geq k\rho$ by mathematical induction: For the base case $k = 0$, $\theta_0 = 0$, so $\theta_0^\top \theta^* = 0 = 0 \cdot \rho$. Assume that when $k = m$, $\theta_m^\top \theta^* \geq m\rho$. Then when $k = m + 1$,

$$\begin{aligned} \theta_{m+1}^\top \theta^* &\geq \theta_m^\top \theta^* + \rho \\ &\geq m\rho + \rho \\ &= (m + 1)\rho. \end{aligned} \quad (27)$$

By induction, $\theta_k^\top \theta^* \geq k\rho$ for all $k \geq 0$.

(c) According to the expansion formula of the vector norm, taking the square norm of $\theta_k = \theta_{k-1} + y_{k-1}x_{k-1}$, we can obtain

$$\|\theta_k\|^2 = \|\theta_{k-1}\|^2 + 2y_{k-1}(\theta_{k-1}^\top x_{k-1}) + \|y_{k-1}x_{k-1}\|^2. \quad (28)$$

Since $y_{k-1} \in \{+1, -1\}$, then $y_{k-1} = 1$. Further, we can obtain $\|y_{k-1}x_{k-1}\|^2 = y_{k-1}^2 \|x_{k-1}\|^2 = \|x_{k-1}\|^2$. Meanwhile, the misclassifications of (x_{k-1}, y_{k-1}) by θ_{k-1} means that the sign of $\theta_{k-1}^\top x_{k-1}$ and y_{k-1} are opposite, i.e., $y_{k-1}(\theta_{k-1}^\top x_{k-1}) < 0$.

Substituting these two results into the norm expansion formula, we can get

$$\|\theta_k\|^2 \leq \|\theta_{k-1}\|^2 + 2 \cdot 0 + \|x_{k-1}\|^2 = \|\theta_{k-1}\|^2 + \|x_{k-1}\|^2. \quad (29)$$

(d) From the conclusion of (c), there is $\|\theta_k\|^2 \leq \|\theta_{k-1}\|^2 + \|x_{k-1}\|^2$. Since R is the maximum norm of all samples x_i , we have

$$\|\theta_k\|^2 \leq \|\theta_{k-1}\|^2 + R^2. \quad (30)$$

Recursively expanding this inequality for $\theta_{k-1}, \theta_{k-2}, \dots, \theta_0$ in turn, we can get

$$\|\theta_k\|^2 \leq \|\theta_0\|^2 + \sum_{i=0}^{k-1} \|x_i\|^2. \quad (31)$$

Since the initial parameter $\theta_0 = 0$, then $\|\theta_0\|^2 = 0$. The summation term $\sum_{i=0}^{k-1} \|x_i\|^2$ contain k term, and each term does not exceed R^2 . Therefore, we can obtain

$$\sum_{i=0}^{k-1} \|x_i\|^2 \leq kR^2. \quad (32)$$

Finally, we have

$$\|\theta_k\|^2 \leq kR^2. \quad (33)$$

(e) Using $\theta_k^\top \theta^* \geq k\rho$ in (b) and $\|\theta_k\|^2 \leq kR^2$ in (d), since all terms are positive, we can obtain that

$$\frac{\theta_k^\top \theta^*}{\|\theta_k\|} \geq \frac{k\rho}{\sqrt{k}R} = \sqrt{k} \frac{\rho}{R} \quad (34)$$

Then, with the Cauchy-Schwarz inequality to bound the number of iterations, we have $\theta_k^\top \theta^* \leq \|\theta_k\| \|\theta^*\|$, i.e., $\frac{\theta_k^\top \theta^*}{\|\theta_k\|} \leq \|\theta^*\|$. Combining $\sqrt{k} \frac{\rho}{R} \leq \frac{\theta_k^\top \theta^*}{\|\theta_k\|} \leq \|\theta^*\|$, we get

$$k \leq \frac{R^2 \|\theta^*\|^2}{\rho^2}, \quad (35)$$

which shows that the perceptron cannot iterate infinitely. Let $\bar{k} = \left\lfloor \frac{R^2 \|\theta^*\|^2}{\rho^2} \right\rfloor$ then after at most \bar{k} iterations, there will be no more misclassified samples and the algorithm terminates.

Problem 5. Pocket Algorithm for Non-Separable data

Solution

(1) The implementation process is in the code file p5.

(2)

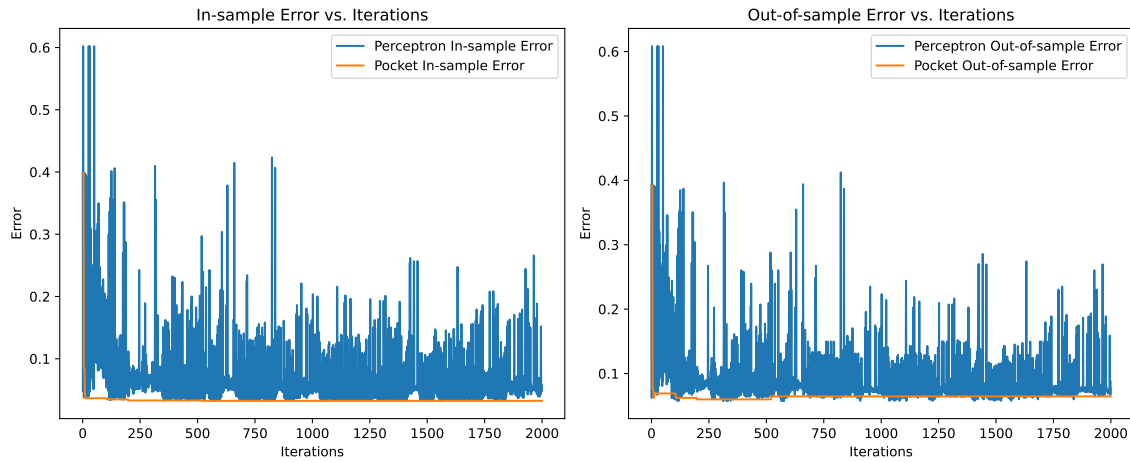


Figure 2: In-sample and Out-of-sample Error Comparison: Perceptron vs. Pocket Algorithm

(3)




Figure 3: Decision Boundaries for Perceptron and Pocket Algorithms on Classifying Digits '1' and '6'

Code of P3

```
1 import numpy as np
2 import matplotlib.pyplot as plt
3
4 d = 50 #feature dimension
5
6 # Load the dataset
7 X = np.load('p3/data/X.npy')
8 y = np.load('p3/data/y.npy')
9 print ("data shape: ", X.shape, y.shape)
10
11 theta_star = np.load('p3/data/theta_star.npy')
12
13 ##### part (1): least square estimator #####
14
15 # Calculate the least square solution
16 theta_hat = np.linalg.inv(X.T @ X) @ X.T @ y
17
18 Error_LS = np.linalg.norm(theta_hat - theta_star, 2)
19 print('Estimator approximated by LS:',Error_LS)
20
21 ##### part (2): L1 estimator #####
22 mu = 1e-5 # smoothing parameter
23 alpha = 0.001 # stepsize
24 T = 1000 # iteration number
25
26 # random initialization
27 theta = np.random.randn(d,1)
28
29 Error_huber = []
30
31 for _ in range(1, T):
32
33     # calculate the l2 error of the current iteration
34     Error_huber.append(np.linalg.norm(theta-theta_star, 2))
35
36     # Calculate the residual
37     r = y - X @ theta
38
39     # Calculate the components of the gradient based on Huber loss
40     g = np.where(np.abs(r) <= mu, r / mu, np.sign(r))
41
42     # Calculate the final gradient
43     grad = -X.T @ g
44
45     #gradient descent update
46     theta = theta - alpha * grad
47
```

```
48 ##### plot the figure #####
49 plt.figure(figsize=(10,5))
50 plt.yscale('log',base=2)
51 plt.plot(Error_huber, 'b-')
52 plt.title(r'$\ell_1$ estimator approximated by Huber')
53 plt.ylabel(r'$||\theta - \theta^*||_2$')
54 plt.xlabel('Iteration')
55 # plt.grid(True)
56 plt.show()
```

Code of P5

```
1     import scipy.io
2 import matplotlib.pyplot as plt
3 import numpy as np
4
5 def load_mat(path, d=16):
6     data = scipy.io.loadmat(path)['zip']
7     size = data.shape[0]
8     y = data[:, 0].astype('int')
9     X = data[:, 1:].reshape(size, d, d)
10    return X, y
11
12 def cal_intensity(X):
13     """
14     X: (n, d, d), input data
15     return intensity: (n, 1)
16     """
17     n = X.shape[0]
18     return np.mean(X.reshape(n, -1), 1, keepdims=True)
19
20 def cal_symmetry(X):
21     """
22     X: (n, d, d), input data
23     return symmetry: (n, 1)
24     """
25     n, d = X.shape[:2]
26     Xl = X[:, :, :int(d/2)]
27     Xr = np.flip(X[:, :, int(d/2):], -1)
28     abs_diff = np.abs(Xl-Xr)
29     return np.mean(abs_diff.reshape(n, -1), 1, keepdims=True)
30
31 def cal_feature(data):
32     intensity = cal_intensity(data)
33     symmetry = cal_symmetry(data)
34     feat = np.hstack([intensity, symmetry])
35     return feat
36
```

```

37 def cal_feature_cls(data, label, cls_A=1, cls_B=5):
38     """ calculate the intensity and symmetry feature of given classes
39     Input:
40         data: (n, d1, d2), the image data matrix
41         label: (n, ), corresponding label
42         cls_A: int, the first digit class
43         cls_B: int, the second digit class
44     Output:
45         X: (n', 2), the intensity and symmetry feature corresponding to
46             class A and class B, where n' = cls_A# + cls_B#.
47         y: (n', ), the corresponding label {-1, 1}. 1 stands for class A,
48             -1 stands for class B.
49     """
50     feat = cal_feature(data)
51     indices = (label==cls_A) + (label==cls_B)
52     X, y = feat[indices], label[indices]
53     ind_A, ind_B = y==cls_A, y==cls_B
54     y[ind_A] = 1
55     y[ind_B] = -1
56     return X, y
57
58 def plot_feature(feature, y, plot_num, ax=None, classes=np.arange(10)):
59     """plot the feature of different classes
60     Input:
61         feature: (n, 2), the feature matrix.
62         y: (n, ) corresponding label.
63         plot_num: int, number of samples for each class to be plotted.
64         ax: matplotlib.axes.Axes, the axes to be plotted on.
65         classes: array(0-9), classes to be plotted.
66     Output:
67         ax: matplotlib.axes.Axes, plotted axes.
68     """
69     cls_features = [feature[y==i] for i in classes]
70     marks = ['s', 'o', 'D', 'v', 'p', 'h', '+', 'x', '<', '>']
71     colors = ['r', 'g', 'b', 'c', 'm', 'y', 'k', 'cyan', 'orange', 'purple']
72     if ax is None:
73         _, ax = plt.subplots()
74     for i, feat in zip(classes, cls_features):
75         ax.scatter(*feat[:plot_num].T, marker=marks[i], color=colors[i], label=str(i))
76     plt.legend(loc='upper right')
77     plt.xlabel('intensity')
78     plt.ylabel('asymmetry')
79     return ax
80
81 def cal_error(theta, X, y, thres=1e-4):
82     """calculate the binary error of the model w given data (X, y)
83     theta: (d+1, 1), the weight vector
84     X: (n, d), the data matrix [X, y]
85     y: (n, ), the corresponding label

```

```
86     """
87     # Add a bias term to X
88     X_b = np.hstack([np.ones((X.shape[0], 1)), X])
89     out = X_b @ theta - thres
90     pred = np.sign(out)
91     err = np.mean(pred.squeeze() != y)
92     return err
93
94 # prepare data
95 train_data, train_label = load_mat('p5/train_data.mat') # train_data: (7291, 16, 16),
96                   train_label: (7291, )
97 test_data, test_label = load_mat('p5/test_data.mat') # test_data: (2007, 16, 16),
98                   train_label: (2007, )
99
100 cls_A, cls_B = 1, 6
101 X, y, = cal_feature_cls(train_data, train_label, cls_A=cls_A, cls_B=cls_B)
102 X_test, y_test = cal_feature_cls(test_data, test_label, cls_A=cls_A, cls_B=cls_B)
103
104 # Add a bias term to the feature matrices
105 X_b = np.hstack([np.ones((X.shape[0], 1)), X])
106 X_test_b = np.hstack([np.ones((X_test.shape[0], 1)), X_test])
107
108 # train
109 iters = 2000
110 d = 2
111 num_sample = X.shape[0]
112 threshold = 1e-4
113
114 # Perceptron and Pocket algorithm initialization
115 theta_p = np.zeros((d + 1, 1))
116 theta_pocket = np.zeros((d + 1, 1))
117 best_theta = np.zeros((d + 1, 1))
118 min_err = cal_error(best_theta, X, y)
119
120 # Lists to store errors
121 err_in_p = []
122 err_out_p = []
123 err_in_pocket = []
124 err_out_pocket = []
125
126 for iterate in range(iters):
127     # Perceptron
128     pred = np.sign(X_b @ theta_p)
129     misclassified_indices = np.where(pred.squeeze() != y)[0]
130
131     if len(misclassified_indices) > 0:
132         # Pick a random misclassified point
133         random_index = np.random.choice(misclassified_indices)
134         xi = X_b[random_index, :].reshape(-1, 1)
```

```
133     yi = y[random_index]
134     theta_p = theta_p + yi * xi
135
136     # Pocket
137     pred_pocket = np.sign(X_b @ theta_pocket)
138     misclassified_indices_pocket = np.where(pred_pocket.squeeze() != y)[0]
139
140     if len(misclassified_indices_pocket) > 0:
141         # Pick a random misclassified point
142         random_index_pocket = np.random.choice(misclassified_indices_pocket)
143         xi_pocket = X_b[random_index_pocket, :].reshape(-1, 1)
144         yi_pocket = y[random_index_pocket]
145         theta_pocket = theta_pocket + yi_pocket * xi_pocket
146
147         # Check if the new theta is better
148         current_err = cal_error(theta_pocket, X, y)
149         if current_err < min_err:
150             min_err = current_err
151             best_theta = theta_pocket.copy()
152
153     # Calculate and store errors
154     err_in_p.append(cal_error(theta_p, X, y))
155     err_out_p.append(cal_error(theta_p, X_test, y_test))
156     err_in_pocket.append(cal_error(best_theta, X, y))
157     err_out_pocket.append(cal_error(best_theta, X_test, y_test))
158
159
160     # plot Er_in and Er_out
161     plt.figure(figsize=(12, 5))
162     plt.subplot(1, 2, 1)
163     plt.plot(range(iters), err_in_p, label='Perceptron In-sample Error')
164     plt.plot(range(iters), err_in_pocket, label='Pocket In-sample Error')
165     plt.xlabel('Iterations')
166     plt.ylabel('Error')
167     plt.title('In-sample Error vs. Iterations')
168     plt.legend()
169
170     plt.subplot(1, 2, 2)
171     plt.plot(range(iters), err_out_p, label='Perceptron Out-of-sample Error')
172     plt.plot(range(iters), err_out_pocket, label='Pocket Out-of-sample Error')
173     plt.xlabel('Iterations')
174     plt.ylabel('Error')
175     plt.title('Out-of-sample Error vs. Iterations')
176     plt.legend()
177     plt.tight_layout()
178     plt.show()
179
180
181     # plot decision boundary
```

```
182 def plot_decision_boundary(X, y, theta, ax, title):
183     # Plot data points
184     ax.scatter(X[y==1][:, 0], X[y==1][:, 1], marker='+', label='1')
185     ax.scatter(X[y==-1][:, 0], X[y==-1][:, 1], marker='*', label='6')
186
187     # Plot decision boundary
188     x1_min, x1_max = ax.get_xlim()
189     x1 = np.array([x1_min, x1_max])
190
191     #  $w_0 + w_1x_1 + w_2x_2 = 0 \Rightarrow x_2 = (-w_0 - w_1x_1) / w_2$ 
192     w = theta.squeeze()
193     if w[2] != 0:
194         x2 = (-w[0] - w[1] * x1) / w[2]
195         ax.plot(x1, x2, 'y-', label='Decision Boundary')
196
197     ax.set_xlabel('Intensity')
198     ax.set_ylabel('Asymmetry')
199     ax.set_title(title)
200     ax.legend()
201
202
203 fig, axes = plt.subplots(1, 2, figsize=(12, 5))
204
205 # Plot for 500 data points
206 plot_num = 500
207 indices_1 = np.where(y == 1)[0][:plot_num]
208 indices_6 = np.where(y == -1)[0][:plot_num]
209 plot_indices = np.concatenate([indices_1, indices_6])
210 X_plot, y_plot = X[plot_indices], y[plot_indices]
211
212 plot_decision_boundary(X_plot, y_plot, theta_p, axes[0], 'Perceptron Decision Boundary')
213 plot_decision_boundary(X_plot, y_plot, best_theta, axes[1], 'Pocket Decision Boundary')
214
215 plt.tight_layout()
216 plt.show()
```