ASSIGNMENT 2: PROMPT ENGINEERING & AGENT

The Teaching Team of Natural Language Processing
The Chinese University of Hong Kong, Shenzhen
Shenzhen, China
`shunianchen@link.cuhk.edu.cn`

You could modify any format as you want and add anything if you think it is necessary. Be concise and say something that matters.

1 INTRODUCTION

Prompt engineering is a relatively new discipline for developing and optimizing prompts to efficiently use language models (LMs) for a wide variety of applications and research topics. Prompt engineering skills help to better understand the capabilities and limitations of large language models (LLMs). Researchers use prompt engineering to improve the capacity of LLMs on a wide range of common and complex tasks such as question answering and arithmetic reasoning. Developers use prompt engineering to design robust and effective prompting techniques that interface with LLMs and other tools. Additionally, prompt engineering is essential for creating AI agents - autonomous systems that can perceive their environment, make decisions, and take actions to achieve specific goals. By carefully crafting prompts, engineers can guide these agents to perform complex, multi-step tasks, reason about their actions, and interact with both users and other digital tools in more sophisticated ways.

2 PRELIMINARY

We have prepared a Colab code `https://colab.research.google.com/drive/1P57pWKROW7zcnXImc0ajVe8OKfiEzh_Z?usp=sharing` for you. This code will guide you through the following:

- How to efficiently call the DeepSeek API.
- Useful Prompt Engineering techniques.
- How to get started with Task 1 and Task 2.

We strongly recommend that you take a quick look at the code first.

3 TASKS

You are encouraged to complete **one and only one task** listed below. To achieve a high score, you must carefully refine the prompt to enhance task performance. You need to implement and compare the results from **at least 3 prompting techniques** introduced in the tutorial or from the `https://www.promptingguide.ai/techniques`, and **an agent-based method**. Additionally, you must include a comparative analysis between different methods, evaluating how LLM agents perform relative to traditional prompting methods across metrics such as accuracy, completion time, reasoning quality, and task complexity handling.

3.1 TASK 1: LLMS AS A KNOWLEDGEABLE DOCTOR

The pharmacist licensure exam is a cornerstone in the pharmacy profession, ensuring that candidates possess the requisite knowledge and skills for safe and effective practice. Its significance lies not only in validating credentials but also in safeguarding public health, enabling professional recognition, and ensuring adherence to legal and regulatory standards.

An example prompt for multi-choice questions

Prompt:
你是一个药剂师考试能手，每次都考100分，这道题对你来说不在话下，深呼吸，并一步一步思考，并给出正确的答案。回答格式为答案是A、B、C或者D。【最佳选择题】根据健康中国战略，推进健康中国建设主要遵循的原则不包括哪个选项？
A: "健康优先"，
B: "改革创新"，
C: "科学发展"，
D: "公开透明"

- -

Expected Output:
根据健康中国战略，推进健康中国建设主要遵循的原则不包括选项D: "公开透明"。

Figure 1: An example prompt for multi-choice questions.

Advanced models like DeepSeek have significant potential in exam preparation, boasting an extensive knowledge base and the capability to provide in-depth explanations and clarify complex concepts. However, despite the prowess of such large models, if prompts are not designed appropriately, the information retrieved might be inaccurate or incomplete, potentially hindering success in the pharmacist exam.

Input: Given Multi-choice questions.

Output: Select the correct answer.

Example: Refer to 3.1.

Criteria: Better accuracy.

Code and Data:

- Data description: 100 questions sampled from 2021 National Pharmacist Professional Qualification Examination real questions
- Code and data: `https://github.com/NLP-Course-CUHKSZ/NLP-course-cuhksz.github.io/tree/main/Assignments/Assignment2/task1`

3.2   TASK 2: LLMS FOR AI FEEDBACK

The final stage of large language model training involves reinforcement learning through feedback. Such feedback can come from either human experts Ouyang et al. (2022) or AI Bai et al. (2022). This feedback is used to learn a reward model, with data defined in a triplet form. This triplet comes from a question, two answers, and a choice by a human or AI on which answer is better.

The triplet consists of three elements: a `question`, the `chosen answer`, and the `rejected answer`. You are asked to use DeepSeek to provide the feedback, namely, choose the preferred one. Note that the feedback is highly biased by the order of placed answers, please shuffle the order of answers when using DeepSeek for preference feedback.

Input: A question Q and two answers A1 and A2.

Output: Which one is chosen and rejected.

Example: Refer to 3.2.

Criteria: The more correlated to the ground truth, the better.

Code and Data:

An example prompt for AI feedback

Prompt:
Given a questions and two answers. You are a smart guy and please tell which answer better answers the given question.
Questions: Do you know why turkeys became the official food of thanksgiving?

Answer1: To be honest, I don't know anything about that. I know that I'm meant to know a lot about history and current events, but I haven't been programmed with those particular facts, sorry.

Answer2: I know that the American Indians were the first to celebrate the first thanksgiving, and the first European settlers came to New England in the early 1600s, and it is likely that they celebrated the first thanksgiving in the late 1600s. However, it is also likely that some European settlers on the continent celebrated the first thanksgiving in the mid-1500s. A lot of people think that the main factor in thanksgiving is that the settlers in the new world were thankful for finding a new land, and that turkey was their traditional food. Another factor that has often been thought to be important is the Pilgrims' relationship with the Native Americans, and the feast may be one way of trying to show them respect. I can't tell you the definitive answer, but maybe this will help you figure it out?

- -

Expected Output:
Answer2 better answers the given question.

Answer1 states that the AI doesn't have knowledge about why turkeys became the official food of Thanksgiving, so it doesn't provide any relevant information.

Answer2, on the other hand, provides historical context and possible reasons for why turkeys became the traditional food of Thanksgiving. It mentions the first celebrations of Thanksgiving by American Indians and European settlers, the settlers' gratitude for finding a new land, and the Pilgrims' relationship with Native Americans. While it doesn't provide a definitive answer, it offers insights and possibilities that can help in understanding the origins of the tradition.

Figure 2: An example prompt for AI feedback.

- Data description: 100 answer pairs sampled from the most downloaded RLHF dataset of HuggingFace `https://huggingface.co/datasets/Anthropic/hh-rlhf`

- Code and data: `https://github.com/NLP-Course-CUHKSZ/ NLP-course-cuhksz.github.io/tree/main/Assignments/ Assignment2/task2`

--

3.3 TASK 3: LLMS AS A DATA QUALITY EVALUATOR

As we know in data science, 'Garbage in, garbage out'; Data is always the most focused part in developing large language models. One could use large language models to evaluate data quality and select better-quality data for training.

Input: A dataset $\mathcal{D}$

Output: A cleaner or better-quality dataset $\mathcal{D}'$

Criteria: You could define the criteria

Data examples: you could filter data in terms of quality for one of the following datasets Peng et al. (2023).

- English Alpaca [1]
- Chinese Alpaca [2]

3.4 TASK 4: ANY OTHER TOPICS

You could choose any other topics you like if it could test prompt engineering.

4 SUBMISSION FORMAT

You are required to submit a short report using the template on `https://www.overleaf.com/read/jpcrtgzjjdry#36db1a`. The report could include the following sections:

- Code to call DeepSeek API if it has coding part.
- Some prompts and their design philosophy.
- How agent is utilized to solve the task
- Some (e.g., two) cases.
- Illustrate and compare the results if it has some.
- Some insights learned (optional).

Ensure that the report is submitted as a PDF via the BB system. The deadline for Assignment 2 is **11:59 P.M. October 31, 2025**.

5 EVALUATION CRITERIA

The Evaluation criteria of this assignment is designed as below:

- **[25 marks]** The quality of prompting engineering. Sorry, this is a black box.
- **[75 marks]** The quality of report. Among these 75 marks, 25 masks are for the clarity of report, 25 marks are for the experimental results (e.g., completeness, experimental rigor, and technical excitement), and 25 marks are for experimental analysis (including case study, sensitiveness of each individual experiment component, ablation study and any other interesting discussions.)
- **[5 mark Bonus]** It is given when there is significance of the selected task if you do not pick any of the tasks in Sec. 3. A good research/engineering taste is always encouraged. This is not guaranteed if you pick a new task, we have to evaluate the significance by a subjective point of view.
- **[5 mark Bonus]** This is for clean and clear codes if it has. This is not guaranteed if you attach you code as supplementary materials.

Note that the marks you can get from this assignment are capped at 100.

---

[1] `https://github.com/Instruction-Tuning-with-GPT-4/GPT-4-LLM/blob/main/data/alpaca_gpt4_da`

[2] `https://github.com/Instruction-Tuning-with-GPT-4/GPT-4-LLM/blob/main/data/alpaca_gpt4_data_zh.json`

ACKNOWLEDGMENT

REFERENCES

Yuntao Bai, Saurav Kadavath, Sandipan Kundu, Amanda Askell, Jackson Kernion, Andy Jones, Anna Chen, Anna Goldie, Azalia Mirhoseini, Cameron McKinnon, et al. Constitutional ai: Harmlessness from ai feedback. *arXiv preprint arXiv:2212.08073*, 2022.

Long Ouyang, Jeffrey Wu, Xu Jiang, Diogo Almeida, Carroll Wainwright, Pamela Mishkin, Chong Zhang, Sandhini Agarwal, Katarina Slama, Alex Ray, et al. Training language models to follow instructions with human feedback. *Advances in Neural Information Processing Systems*, 35: 27730–27744, 2022.

Baolin Peng, Chunyuan Li, Pengcheng He, Michel Galley, and Jianfeng Gao. Instruction tuning with gpt-4, 2023. URL `https://arxiv.org/abs/2304.03277`.